

1 b) [0.5 points]

We define a function that takes μ_1 , μ_2 and Σ as arguments and computes $p(1|2)$.

```
miscl_theo <- function(mean1, mean2, sigma){  
  mah_dist <- sqrt( t(mean1 - mean2) %*% solve(sigma) %*% (mean1-mean2))  
  pnorm( -0.5* mah_dist)  
}
```

We can now compute the theoretical probability of misclassification for the given values of μ_1, μ_2 and Σ .

```
mu1 <- c(1, -2)  
mu2 <- c(2,0)  
covSigma <- diag(c(1,1))  
  
miscl_theo(mean1 = mu1, mean2 = mu2, sigma = covSigma)
```

```
##           [,1]  
## [1,] 0.1317762
```

Thus, the probability of making a wrong classification is 13.1776239%.

We draw a sample from the corresponding Gaussian mixture model.

```
samp1 <- mvtnorm::rmvnorm(100, mean = mu1, sigma = diag(c(1,1)))  
samp2 <- mvtnorm::rmvnorm(100, mean = mu2, sigma = diag(c(1,1)))  
  
gmm_sample <- data.frame(rbind(samp1, samp2) , Class = rep(1:2, each = 100))
```

Since the components are independent here (covariance is 0, for normal variables that coincides with independence) one doesn't need the `mvtnorm` package. One can also use base R's `rnorm()`, which samples from the univariate normal distribution, as follows. First, for each component, sample 200 observations from $\mathcal{N}(0,1)$:

```
gmm_sample <- data.frame( X1 = rnorm(200), X2 = rnorm(200), Class = rep(1:2, each = 100))
```

The joint distribution is then $\mathcal{N}(0, I_2)$ with I_2 the identity matrix on \mathbb{R}^2 .

Now we adjust the means (note that for $X \sim \mathcal{N}(0,1)$ one has $\mu + X \sim \mathcal{N}(\mu, 1)$).

```
gmm_sample[1:100, 1] <- gmm_sample[1:100, 1] + 1  
gmm_sample[1:100, 2] <- gmm_sample[1:100, 2] - 2  
gmm_sample[101:200, 1] <- gmm_sample[101:200, 1] + 2
```

Now we write a function that estimates the out-of-sample error based on leaving-one-out cross-validation (LOOCV). The function's argument `sample_data` needs to have a column named `Class` containing the class labels.

```
compute_loocv <- function(j, sample_data = gmm_sample){  
  lda_loo <- MASS::lda(Class ~ ., data = sample_data[-j, ]) # fit LDA to all but j-th observations  
  # return 1 (FALSE) if wrong class is predicted, else return 0 (TRUE)  
  predict(lda_loo, newdata = sample_data[j, ])$class != sample_data$Class[j]  
}
```

Then we apply the function for each j in $1:200$ (i.e. each row is left out once), and then we average

```
mean(sapply( 1:200 , compute_loocv))
```

```
## [1] 0.12
```

The LOOCV misclassification rate is 0.12.

c) [0.5 points]

We write a function that performs one iteration, i.e. sampling and computing LOOCV error. The arguments of the function are `seed`, which is set at the beginning of each iteration to make things reproducible, the mean vectors `mean1` and `mean2` as well as the covariance matrix. The default value of the latter is the identity matrix on \mathbb{R}^2 .

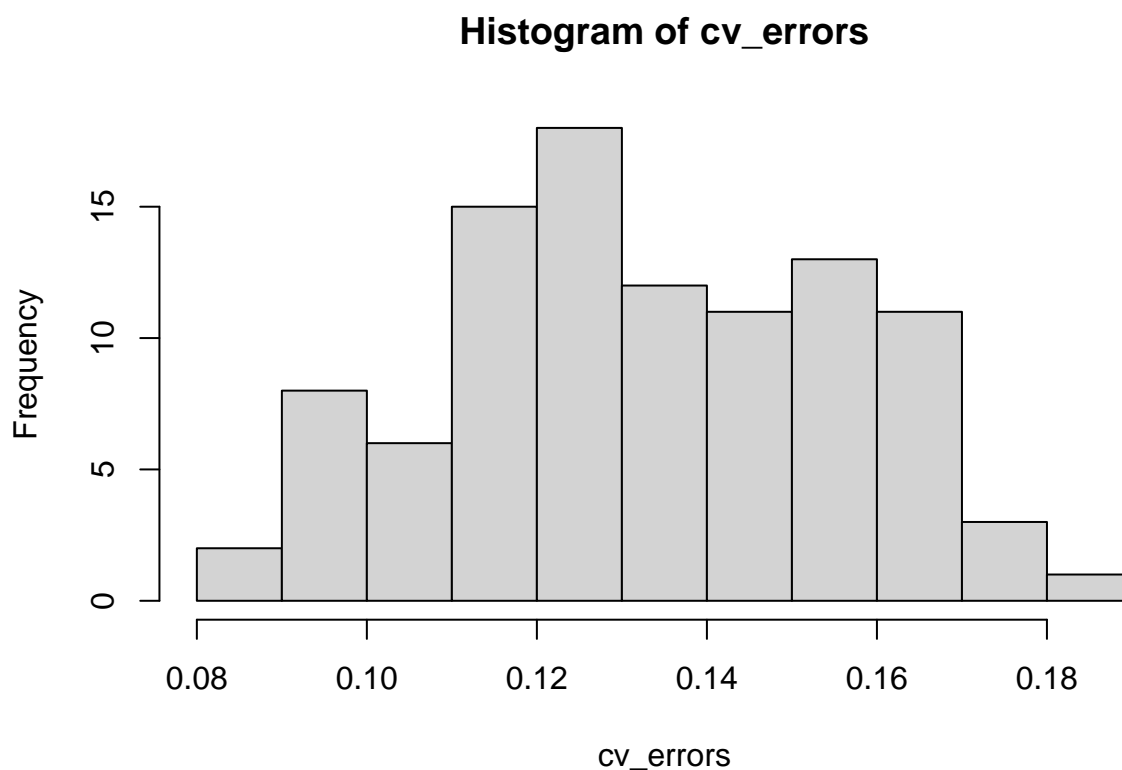
```
iteration <- function(seed , mean1 , mean2, sigma = diag(c(1,1))){
  set.seed(seed)
  samp1 <- mvtnorm::rmvnorm(100, mean = mu1, sigma = sigma )
  samp2 <- mvtnorm::rmvnorm(100, mean = mu2, sigma = sigma)

  gmm_samp <- data.frame(rbind(samp1, samp2) , Class = rep(1:2, each = 100))

  mean(sapply( 1:200 , compute_loocv, sample_data = gmm_samp)) # this is what the function returns
}
```

Then we make 100 iterations, i.e. we apply the above functions with seed from 1 to 100 and make a histogram, compute mean and standard deviation.

```
cv_errors <- sapply(1:100, function(x){iteration(seed = x, mean1 = mu1,
                                                mean2 = mu2, sigma = covSigma)})
hist(cv_errors)
```



```
mean(cv_errors)
```

```
## [1] 0.13545
```

```
sd(cv_errors)
```

```
## [1] 0.02371011
```

```
misc1_theo(mean1 = mu1, mean2 = mu2, sigma = covSigma)
```

```
##           [,1]
```

```
## [1,] 0.1317762
```

We see that the mean is pretty close (deviates in third decimal place) and that the standard deviation is small.

Exercise 2

First of all, we load the required packages and the `penguins` data.

In the text it says to only keep columns containing measurements of quantitative variables, as well as the column containing the penguin's species. Later we will need the column `sex`, too (I forgot to mention it here, sorry). Further, we want to keep only the rows with complete observations on these variables. We generate a new dataframe following these directions.

```
pengus <- dplyr::select(penguins, -c("year", "island"))  
# only complete cases  
pengus <- pengus[complete.cases(pengus), ]
```

a) [0.5 points]

Now we randomly sample 80% of the data as training data. Then we look at the proportions of the different species within the training data.

```
set.seed(4017)  
# sample indices of rows that make up the training data  
ind <- sample(nrow(pengus), nrow(pengus)*.8 )  
  
# training data  
trainsamp <- pengus[ind , ]  
# distribution of species within training data  
table(trainsamp$species)/nrow(trainsamp)
```

```
##
```

```
##      Adelie Chinstrap   Gentoo
```

```
## 0.4774436 0.1729323 0.3496241
```

The test data consists of the observations that are not part of the training data. We compute the distribution of species within the test data.

```
testsamp <- pengus[-ind, ]
```

```
table(testsamp$species)/nrow(testsamp)
```

```
##
```

```
##      Adelie Chinstrap   Gentoo
```

```
## 0.2835821 0.3283582 0.3880597
```

We see that the distributions of the species differ quite drastically for the two datasets. This can affect the model's performance, when the true distribution of classes is not represented well in the training data. Since random splitting into training and test data doesn't guarantee the latter, one should be careful when only using one train-test-split.

We compare to the distribution of classes for the whole dataset.

```
table(pengus$species)/nrow(pengus)
```

```
##
##      Adelie Chinstrap   Gentoo
## 0.4384384 0.2042042 0.3573574
```

b) [0.5 points]

There are many possible ways, and of course there are also packages that provide functions for this kind of sampling. Here, we present one solution that uses some self-written functions.

First, we split the dataframe into a list of three dataframes, each containing the measurements belonging to one species.

```
pengus_split <- split(pengus, pengus$species)
pengus_split
```

```
## $Adelie
## # A tibble: 146 x 6
##   species bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>      <dbl>        <dbl>          <int>        <int> <fct>
## 1 Adelie      39.1          18.7            181         3750 male
## 2 Adelie      39.5          17.4            186         3800 female
## 3 Adelie      40.3          18             195         3250 female
## 4 Adelie      36.7          19.3            193         3450 female
## 5 Adelie      39.3          20.6            190         3650 male
## 6 Adelie      38.9          17.8            181         3625 female
## 7 Adelie      39.2          19.6            195         4675 male
## 8 Adelie      41.1          17.6            182         3200 female
## 9 Adelie      38.6          21.2            191         3800 male
## 10 Adelie     34.6          21.1            198         4400 male
## # ... with 136 more rows
##
## $Chinstrap
## # A tibble: 68 x 6
##   species   bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>        <dbl>        <dbl>          <int>        <int> <fct>
## 1 Chinstrap    46.5          17.9            192         3500 female
## 2 Chinstrap     50           19.5            196         3900 male
## 3 Chinstrap    51.3          19.2            193         3650 male
## 4 Chinstrap    45.4          18.7            188         3525 female
## 5 Chinstrap    52.7          19.8            197         3725 male
## 6 Chinstrap    45.2          17.8            198         3950 female
## 7 Chinstrap    46.1          18.2            178         3250 female
## 8 Chinstrap    51.3          18.2            197         3750 male
## 9 Chinstrap     46           18.9            195         4150 female
## 10 Chinstrap    51.3          19.9            198         3700 male
## # ... with 58 more rows
##
## $Gentoo
## # A tibble: 119 x 6
##   species bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>      <dbl>        <dbl>          <int>        <int> <fct>
## 1 Gentoo     46.1          13.2            211         4500 female
```

```
## 2 Gentoo          50          16.3          230          5700 male
## 3 Gentoo          48.7         14.1          210          4450 female
## 4 Gentoo          50          15.2          218          5700 male
## 5 Gentoo          47.6         14.5          215          5400 male
## 6 Gentoo          46.5         13.5          210          4550 female
## 7 Gentoo          45.4         14.6          211          4800 female
## 8 Gentoo          46.7         15.3          219          5200 male
## 9 Gentoo          43.3         13.4          209          4400 female
## 10 Gentoo         46.8         15.4          215          5150 male
## # ... with 109 more rows
```

Now we write a function that samples a proportion of `train_prop` from the `1:nrow(data)`:

```
get_train_index <- function(data, train_prop = 0.8, seed = 1){
  set.seed(seed) # set seed to make reproducible
  ndata <- nrow(data)
  sample(ndata, floor(ndata*train_prop))
}
```

The function that returns training data based on rows sampled with `get_train_index()`:

```
get_train_data <- function(data, train_prop = 0.8, seed = 1){
  samp_ind <- get_train_index(data = data, train_prop = train_prop, seed = seed)
  train_data <- data[samp_ind, ]
  return(train_data)
}
```

And function that returns test data based on rows sampled in `get_train_index`. Note that, when applying these functions, the same seed must be used for the test data as for the training data.

```
get_test_data <- function(data, train_prop = 0.8, seed = 1){
  samp_ind <- get_train_index(data = data, train_prop = train_prop, seed = seed)
  test_data <- data[-samp_ind, ]
  return(test_data)
}
```

We use the function `purrr::map_dfr` which maps the function `get_train_data()` to the list elements in `pengus_split` and binds the resulting dataframes by rows (dfr means dataframe rowbind).

```
pengus_train <- purrr::map_dfr(pengus_split, get_train_data)
pengus_test <- purrr::map_dfr(pengus_split, get_test_data)
pengus_train
```

```
## # A tibble: 265 x 6
##   species bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>         <dbl>         <dbl>         <int>         <int> <fct>
## 1 Adelie         45.8           18.9           197           4150 male
## 2 Adelie         38.1           17.6           187           3425 female
## 3 Adelie         36            17.9           190           3450 female
## 4 Adelie         34.4           18.4           184           3325 female
## 5 Adelie         39            17.5           186           3550 female
## 6 Adelie         35.7           18            202           3550 female
## 7 Adelie         35.3           18.9           187           3800 female
## 8 Adelie         45.6           20.3           191           4600 male
## 9 Adelie         42.1           19.1           195           4000 male
## 10 Adelie        39.2           19.6           195           4675 male
## # ... with 255 more rows
```

map_dfr() is kind of advanced, you can also use lapply() and assemble by hand.

```
table(pengus_train$species)/nrow(pengus_train)
```

```
##
##      Adelie Chinstrap      Gentoo
## 0.4377358 0.2037736 0.3584906
```

```
table(pengus_test$species)/nrow(pengus_test)
```

```
##
##      Adelie Chinstrap      Gentoo
## 0.4411765 0.2058824 0.3529412
```

c) [0.5 points]

```
lda_peng <- lda(species ~ . - sex, data = pengus_train) # . - sex uses all columns but `sex`
pred_peng_test <- predict(lda_peng, newdata = pengus_test)
conf_test <- table(pred_peng_test$class, pengus_test$species)
1- sum(diag(conf_test))/nrow(pengus_test)
```

```
## [1] 0.01470588
```

The misclassification error on the test set is 0.0147059.

d) [1 point]

```
cv_kfold <- function(K, data){
  n_data <- nrow(data)

  # K must be an integer between 1 and nrow(data)
  if( !( K == round(K)) | K == 0 | K > n_data){stop( "Choose different value for K.")}

  n_subsamp <- floor(n_data/K)      # length of each of the K subsamples

  data <- data[sample(n_data), ]   # shuffle data

  # assign each observation to one of the K subsamples, column 'K' specifies the subsample
  # if n_data is not a multiple of K, assign the last n_data- K*n_subsamp observations randomly

  if( n_subsamp == n_data/K) {
    data$K <- rep(1:K, each = n_subsamp)
  }
  else{
    data$K <- c(rep(1:K, each = n_subsamp), sample(1:K, (n_data - K*n_subsamp)))
  }

  # initialise vector that stores the misclassification rates
  cv_errors <- numeric(K)

  for( k in 1:K){
    # training data in iteration k
    data_train <- subset(data, !( K == k))   # leave out k-th subset
    data_train <- dplyr::select(data_train, - "K") # get rid of column specifying the subset
  }
}
```

```

# test data in iteration k
data_test <- subset(data, K == k)      # k-th subset
data_test <- dplyr::select(data_test, - "K") # get rid of column specifying the subset

lda_fit <- lda(species ~ .,
               data = data_train)      # fit lda to training data
test_pred <- predict(lda_fit, newdata = data_test) # predict class for test data

conf_mat <- table(test_pred$class, data_test$species) # compute confusion matrix

cv_errors[k] <- 1- sum(diag(conf_mat))/nrow(data_test) # misclassification rate in iteration k
}
# return the average misclassification rate and also the vector of misclassification rates,
# in case one is interested in variance
return(list(avg_misclass = mean(cv_errors), errors = cv_errors))
}

```

We apply the function:

```
cv_kfold(K = 10, data = penguins)
```

```

## $avg_misclass
## [1] 0.006060606
##
## $errors
## [1] 0.00000000 0.00000000 0.00000000 0.03030303 0.00000000 0.00000000
## [7] 0.03030303 0.00000000 0.00000000 0.00000000

```

e) [1.5 points]

First, we subset the measurements of Chinstrap species.

```

penguins_chin <- subset(penguins, species == "Chinstrap")
penguins_chin

```

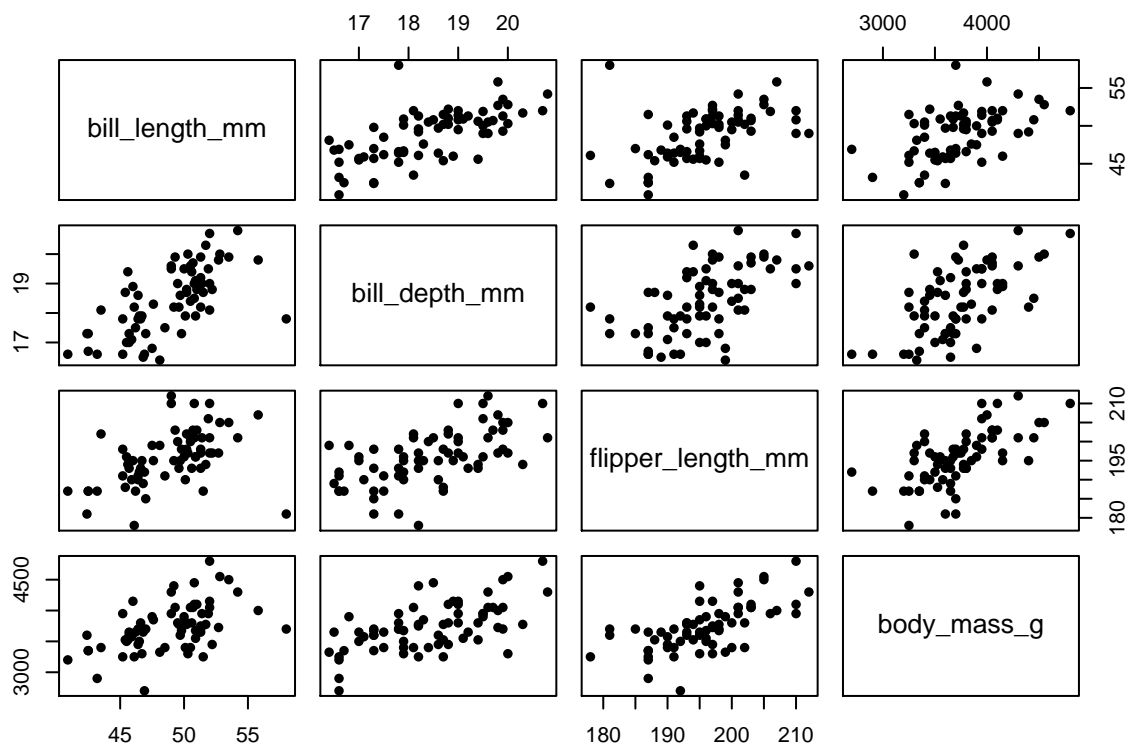
```

## # A tibble: 68 x 6
##   species    bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>          <dbl>          <dbl>             <int>         <int> <fct>
## 1 Chinstrap      46.5            17.9              192          3500 female
## 2 Chinstrap      50              19.5              196          3900 male
## 3 Chinstrap      51.3            19.2              193          3650 male
## 4 Chinstrap      45.4            18.7              188          3525 female
## 5 Chinstrap      52.7            19.8              197          3725 male
## 6 Chinstrap      45.2            17.8              198          3950 female
## 7 Chinstrap      46.1            18.2              178          3250 female
## 8 Chinstrap      51.3            18.2              197          3750 male
## 9 Chinstrap      46              18.9              195          4150 female
## 10 Chinstrap     51.3            19.9              198          3700 male
## # ... with 58 more rows

```

We can make a pairwise scatter plot to check linear relation of predictors and target.

```
plot(penguins_chin[, -c(1,6)], pch = 16)
```



```
# also check correlation matrix
cor(pengus_chin[ , -c(1, 6)])
```

```
##               bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
## bill_length_mm      1.0000000      0.6535362      0.4716073      0.5136383
## bill_depth_mm       0.6535362      1.0000000      0.5801429      0.6044983
## flipper_length_mm   0.4716073      0.5801429      1.0000000      0.6415594
## body_mass_g         0.5136383      0.6044983      0.6415594      1.0000000
```

Assumption of linear relation seems alright.

Now we fit the linear models and estimate the out-of-sample error based on leaving-one-out cv. The Mean squared error is the squared error here, because in each iteration we only have one predicted value.

First model: only bill depth as explanatory variable

```
compute_loocv1 <- function(j){
  lm_loo <- lm( bill_length_mm ~ bill_depth_mm, data = pengus_chin[-j , ]) # leave out j-th observat
  yhat <- predict(lm_loo, newdata = pengus_chin[j, ]) # predict bill_length based on model fit
  (pengus_chin$bill_length_mm[j] - yhat)^2 # return squared error
}
```

Apply function to all indices in 1:nrow(pengus), i.e. leave out each row once, then compute the average.

```
cv_mse1 <- mean(sapply(1:nrow(pengus_chin), compute_loocv1))
cv_mse1
```

```
## [1] 6.638773
```

Second model: body mass and bill depth as explanatory variables.


```
compute_loocv2 <- function(j){
  lm_loo <- lm( bill_length_mm ~ body_mass_g + bill_depth_mm, data = penguins_chin[-j , ]) # leave out
  yhat <- predict(lm_loo, newdata = penguins_chin[j, ]) # predict bill_length based on model fit
  (penguins_chin$bill_length_mm[j] - yhat)^2 # return squared error
}

cv_mse2 <- mean(sapply(1:nrow(penguins_chin), compute_loocv2))
cv_mse2
```

```
## [1] 6.524676
```

Third model: all three variables as predictors

```
compute_loocv3 <- function(j){
  lm_loo <- lm( bill_length_mm ~ flipper_length_mm + body_mass_g + bill_depth_mm ,
              data = penguins_chin[-j , ]) # leave out j-th observation
  yhat <- predict(lm_loo, newdata = penguins_chin[j, ]) # predict bill_length based on model fit
  (penguins_chin$bill_length_mm[j] - yhat)^2 # return squared error
}

cv_mse3 <- mean(sapply(1:nrow(penguins_chin), compute_loocv3))
cv_mse3
```

```
## [1] 7.096585
```

Compare the MSEs:

```
cv_mse2 < cv_mse1
```

```
## [1] TRUE
```

```
cv_mse3 < cv_mse2
```

```
## [1] FALSE
```

```
cv_mse3 < cv_mse1
```

```
## [1] FALSE
```

Model 2 has smallest MSE. However, reduction of MSE from model 2 to model 1 is only marginal. Therefore, one would probably go for model 1 (also ok if you chose model 2).